ONRA

Customer & Lead Management

Product design specification — addendum to PRD Section 6

How the Customer section works: a hybrid model where AI detects what is happening and staff decide what to do about it.

June 2026

Contents

1. Summary

This document specifies how Onra's Customer section works end to end — from a prospect's first signal, through manual sales follow-up, to an AI-managed paying member. It resolves the core tension: when the system should manage a customer's state automatically, and when a human should.

The governing principle: AI owns detection and the default state; staff own the action and the override. No one does data entry, and no one has to remember to look — AI watches and proposes, staff decide and act.

This is the opposite of the incumbents (Mindbody, Momence, Glofox), whose lead management is manual and form-driven. Onra's wedge is the **signal → task engine**: it auto-detects intent and abandonment, then hands staff a ready-made, pre-assigned follow-up task with the reason attached.

> **Why this matters in the UAE market**
> Dubai studios already chase every lead aggressively — it's real, high-touch work. Onra doesn't add a manual pipeline on top; it automates the part they currently do from memory and a WhatsApp list.

2. The core model: one record, two layers

Everything in the Customer section rests on a single idea: a lead and a customer are the same person record at different points in their lifecycle. There is no separate "leads database" and "customers database." There is one record whose state advances over time. Splitting them into two stores creates duplicate profiles and a sync problem — the exact mess studios complain about in their current platforms.

On top of that single record sit two distinct layers that answer two different questions. Keeping them separate is what makes the hybrid coherent instead of contradictory.

2.1 Layer 1 — Lifecycle stage (AI-owned, automatic)

This is the PRD's existing **AI-assigned lifecycle tag** (Section 6.23). It answers ***what is this person actually doing?*** — the objective, behavioural truth, derived continuously from attendance, booking patterns, payment behaviour, and rating history. Staff do not maintain it by hand; the AI does, and it updates on its own as behaviour shifts.

Lifecycle tag	What it means
Lead	Has signalled interest but never converted. No behavioural history yet.
Trialist	Booked or attended a trial / intro offer; not yet on a paid plan.
New Active	Recently converted to a paid plan; habit not yet established.

Lifecycle tag	What it means
Loyal Active	Established, regular, paying member.
At Risk	Attendance or engagement dropping; churn signals present.
Churned	Lapsed; plan expired or cancelled, no recent activity.
Won-back	Previously churned, now re-engaged.
VIP	High value / high engagement; flagged for priority care.

Design rule: staff never drag a person across Layer 1 by hand. It is derived from behaviour. This is the differentiator — protect it. The tag is always shown alongside the reasoning that produced it, so it is transparent rather than a black box.

2.2 Layer 2 — Follow-up status (staff-owned, manual)

This is the thin sales overlay — the Kenko-style stages reviewed. It answers a different question: ***what are we doing about this lead?*** It captures human sales intent that no behavioural signal can know — whether a prospect is worth a call, said "maybe later," or is a dead end.

The critical scoping rule: Layer 2 only exists before conversion. It applies to people whose Layer 1 tag is Lead or Trialist. The moment they buy, the follow-up status retires and Layer 1 takes over completely. This is what stops the two layers from competing — one observes, the other acts, and they live on different parts of the funnel.

Default follow-up statuses (studio-configurable — see Section 4):

Follow-up status	Meaning
New	Just entered; not yet contacted.
Contacted	Reached out; awaiting response.
Trial booked	Has a trial / intro on the calendar.
Follow-up	Needs another touch; warm but not closed.
Won	Converted to a paying member (auto-set on purchase).
Lost	Not interested / went cold. Suppresses re-surfacing.

2.3 How the two layers relate — precedence

Because both layers can speak at once, one precedence rule prevents deadlock:

- **Human intent overrides AI surfacing.** If a staff member marks a lead "Lost / Not interested," the AI stops re-surfacing them for follow-up. The person's stated judgement wins over the system's nudges.

- **Real behaviour overrides everything.** If that same "Lost" lead later books a class, behaviour wins: they re-enter the active lifecycle and the AI takes back over. No one has to remember to undo the "Lost" flag.

In short: AI sets the default, humans can override the default, and real-world behaviour overrides the humans. This ordering is simple to explain to an operator and never traps a record in a wrong state.

3. The signal → task engine (the differentiator)

This is the part that makes Onra better than the incumbents, and the centrepiece of the customer section. Today, studios spot leads to chase from memory. Onra detects the moment automatically and hands staff a finished task.

3.1 How it works

1. **Detect.** Catches intent/abandonment signals automatically
2. **Create the task.** Auto-builds a pre-assigned task with the reason attached. (e.g. "Sara started signup 2 days ago, never booked — call her").
3. **Surface it.** It appears in the Snapshot Insights "Attention needed" widget under "Leads to follow up," ranked by value and freshness.
4. **Act and log.** Staff call and log the outcome (Reached / Follow-up / Not interested). The outcome closes, snoozes, or advances the Layer 2 status.
5. **Feed back.** The outcome updates the lead and the AI; the pipeline self-maintains.

Division of labour: the AI does the watching that humans forget to do; the human does the relationship that AI cannot. Neither does the other's job.

3.2 Trigger events to detect

These are the signals Dubai studios already act on manually. Onra catches them automatically:

Trigger	Auto-generated follow-up
Signup started, not verified / completed	"X began signup and stopped — finish onboarding with them."
App downloaded, no booking within 48h	"X installed the app but hasn't booked — nudge / call."
Enquiry logged (call, walk-in, DM)	"New enquiry from X — follow up."
Trial attended, no rebooking within 7 days	"X did a trial and hasn't returned — win them onto a plan."
First booking cancelled, never rebooked	"X cancelled their first class and went quiet."
Lead form submitted (Meta, web, ClassPass, Gympass)	"New lead from [source] — reach out."

Trigger	Auto-generated follow-up
Prototype validation question During the prototype-validation month, mock this engine and ask the 2–3 pilot studios two things: (1) which of these abandoned-action signals they most want auto-caught, and (2) whether they want a kanban/pipeline view or a simple ranked worklist. Their answers determine how much of Layer 2 to actually build.	

4. Configurable sources and stages

4.1 Lead sources — adopt, and upgrade the PRD

Lead source is attribution: where the person came from. The PRD currently hardcodes this as a fixed list (***referral* | *walk-in* | *social* | *classpass* | *other***). That will be wrong within a quarter — channels change as studios join or drop ClassPass, pick up Gympass, or run Meta lead-form campaigns.

Recommendation: make lead source a studio-editable list in Settings, exactly like the reference app. Seed sensible defaults, let the studio add/rename/remove. This feeds the "Member acquisition source" report (PRD 6.3) and costs almost nothing.

Suggested default sources: Walk-in, Referral, Instagram / Social, Website / Online, Web form, Meta lead form, ClassPass, Gympass, Expired customer, Other.

4.2 Follow-up stages — keep, but scope tightly

Make the Layer 2 follow-up statuses studio-configurable too, but with guardrails so they don't drift back into a heavy manual pipeline that undermines the AI story:

- **Scope to pre-conversion only.** Stages describe un-converted leads. They disappear on purchase.
- **Ship tight defaults.** New, Contacted, Trial booked, Follow-up, Won, Lost. Studios can adjust, but most won't need to.
- **Keep it a status, not a chore.** The engine populates and advances most of it. Manual editing is the exception, not the daily routine.

5. Where it lives in the product

5.1 Leads live in the Customers list, not a separate module

Leads are not a separate section. They are a **segment of the Customers list**, surfaced via stage tabs (All / Leads / Members / Inactive) above the existing table. The same data powers two audiences: the front desk works the full list; Marketing reads the "Leads" segment to build nurture campaigns. One dataset, different lenses — no second database.

Optional pipeline view. For studios that think in pipeline terms, offer a kanban view (columns = follow-up statuses) as an alternative rendering of the same Leads segment.

5.2 What the Customer record carries

Field group	Contents	Owner
Identity	Name, phone (primary), email (optional), DOB, avatar	Captured progressively
Lifecycle tag	Layer 1 AI tag + reason + last-updated	AI
Follow-up status	Layer 2 stage (pre-conversion only)	Staff / engine
Source	Configurable attribution (first-touch)	Captured at entry
Assignment	Owner / assignee, next-follow-up date	Staff / engine
Activity log	Calls, notes, enquiries, outcomes	Staff / engine
Plans & history	Memberships, packages, bookings, payments, ratings	System

6. Designed to scale: boutique today, clubs tomorrow

Onra is boutique-first, but the goal is to also serve larger fitness clubs and gyms over time. The good news: the two-layer model is exactly the right architecture for that, because the difference between a boutique studio and a big club is mostly a difference of degree in Layer 2, not a different system.

Dimension	Boutique studio (now)	Larger club / gym (later)
Sales motion	Light: front desk chases a handful of leads	Heavy: dedicated membership advisors with quotas
Layer 2 depth	Few stages, mostly engine-driven worklist	Full pipeline: routing, assignment rules, SLAs, lead scoring
Lead volume	Low — worklist is enough	High — auto-routing across multiple reps becomes essential
Plan model	Credits / class packs	Contracts, annual & corporate memberships, access control
Roles	Front desk, operator	Add a sales / membership-advisor role

6.1 What to build flexibly now so clubs are cheap later

- **Configurable stages and sources.** Already the plan. A club simply configures a deeper pipeline; the boutique keeps it shallow. No rearchitecture.
- **Don't hardcode "credits-only."** Keep the plan / product model general enough that contract-based and access-based memberships can slot in later without breaking the customer record.
- **Assignment as a first-class field.** Even if boutiques rarely use it, store owner/assignee on the lead now so multi-rep routing is a config change, not a schema change later.
- **Lean on the custom-role builder.** The PRD's role system (Section 3) already lets a club add a membership-advisor role without new code.
- **Keep the signal → task engine generic.** The triggers (abandoned signup, no rebook, enquiry) are universal across boutique and big-box. Build them as configurable rules, not hardcoded boutique assumptions.

The flexibility principle

Build the hybrid so it scales from a light worklist (boutique) to a full sales pipeline (club) by configuration, not rebuild. The boutique sees a simple version; the club turns on depth. Same model, same record, same engine — different settings.